

Single Variable Regression: Introduction

EC420 MSU Spring 2021

Justin Kirkpatrick

Last updated January 19, 2021

Goal:

1. Introduce the problem we'll be working on for a while
2. Define the **Population Regression Function**
3. Intuition of "fitting a line"
4. Define the assumptions for OLS
5. Ordinary Least Squares estimator
6. Computing OLS estimates in the **Sample Regression Function**
7. Descriptive analysis vs. causal
8. Other methods of calculating OLS

We have some data on two (or more, later) variables that we think move together in an interesting way.

- (Insert one of many examples we've talked about before here)


We want to quantify and test this relationship

- Predict a change
- Test a theory
- Win a bet?

We have a **sample**, but want to predict/test something about the population

The problem at hand...

Wage data used in Wooldridge `wage2`




Data from Blackburn and Neumark (1992), "Unobserved Ability, Efficiency Wages, and Interindustry Wage Differentials" *Quarterly Journal of Economics* 107, 1421-1436

The problem at hand

The data looks like this:

wage	educ
769	12
808	18
825	14
650	12
562	11
1400	16



- $N = 935$
- $\overline{wage} = 957.95$
- $\overline{educ} = 13.47$

What we'd like to have is a function that tells us how *wage* and *educ* move together in the **population**

In a perfect world, we would have some function for $X = educ$ and $Y = wage$:

$$g(x) = y$$

Where we give the function any realization of x , and it spits out exactly y .

But that isn't going to happen

Think about the data we just looked at - $educ = 12$ we observed $wage = 769$ and $wage = 650$. The dream function doesn't exist! There are other things not accounted for besides $educ$.

So we settle for something that tells us about the **expectation** of Y . The Population Regression Function

$$E[Y|X] = \beta_0 + \beta_1 X$$

The *Population Regression Function* (PRF) describes the relationship between X and the **conditional expectation** of Y .

- X and Y are random variables
- β_0 and β_1 are **population parameters**
- We have restricted the $E[Y|X]$ to be a *linear* function of X .
 - It can be drawn as a straight line with an intercept and constant slope
 - We will be estimating β_0 and β_1


The PRF:

$$E[Y|X] = \beta_0 + \beta_1 X$$


Let $Y = \text{wage}$ and $X = \text{educ}$

- $E[Y|X = x]$ gives us the expectation of Y (wage) conditional on some realized value of $X = x$ (educ)
- So, if $\text{educ} = 16$, then $E[Y|X = 16] = \beta_0 + \beta_1 \times 16$
 - We can plug in any x_i and get the **expected value** of the paired y_i


Question: Will the PRF return exactly y_i given a value x_i ?



Ch. 2.1 of Wooldridge, example of a conceptual PRF. The line defines the PRF, the expectation of Y conditional on X



This is the wage data. Each "blob" is an empirical histogram of the data for that value of *educ* (they are symmetrical). This is called a *violin plot*. It is the empirical counterpart of the [previous plot from Wooldridge](#)



Each point is the sample mean for each value of *educ*.

A (linear) PRF would be the straight line that best connects the points. **Regression fits that line.** A brief look at the line shows that it certainly won't be perfect!

What happens, then, if we want to write Y exactly?

The PRF gives us the *expectation* of Y

- So we add a **stochastic error term**, the difference between $E[Y|X]$ and Y :

$$Y = E[Y|X] + U = \beta_0 + \beta_1 X + U$$

This is the stochastic population regression function

U is also the **population error term**, and is itself a **random variable**.

- It must be that $E[U] = 0$

Now we can write our **simple linear regression model**:

$$y = \beta_0 + \beta_1 x + u$$

This is a statement about the relationship between observed realizations (y_i, x_i) based on the population parameters β_0, β_1

We will call u the **error term** - it is the difference between the conditional expected mean and the observed y_i given a value of x_i .

- It might be different for two identical realizations of x_i

Naturally, we would think that the "right" value of the population parameters, $\beta = \{\beta_0, \beta_1\}^*$, minimizes all of the u_i values in a sample.

* A parameter vector is just a list of numbers.

The Sample Regression Function

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

The "hats" are important

They mean we have a *sample estimate* of the population parameters.

- β_0, β_1 are the population
- $\hat{\beta}_0, \hat{\beta}_1$ are the sample estimates and will change when the sample changes
 - So they are random variables!

Where did u go?

Since we have a hat on y_i , there is no u , but $\hat{y}_i \neq y_i$.

- Define $\hat{u}_i = \hat{y}_i - y_i$.
- \hat{u}_i is the *residual*.

To summarize:

The *PRF* is

$$E[Y|X] = \beta_0 + \beta_1 X$$

The simple linear regression model is:

$$y = \beta_0 + \beta_1 x + u$$

The SRF is:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

And if we want to write the sample regression model:

$$y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + \hat{u}_i$$

How do we get those $\hat{\beta}$'s in the SRF?

We make two assumptions:

First, if the expectation of Y equals $\beta_0 + \beta_1 X$, then *in expectation*, $E[U] = 0$.
Because:

$$E[Y|X] = \beta_0 + \beta_1 X \quad \text{and} \quad Y = \beta_0 + \beta_1 X + U$$

Second, our first assumption should hold no matter what x is. So, it should be true that $E[U|X] = 0$ for **all** possible values of X .

There are very important assumptions as they will define our Sample Regression Function (SRF).

Let's make these assumptions formal:

1. $E[U] = 0$.

- As long as there is a β_0 (regardless of β_1), this is true. We call this assumption **trivial**.

2. $E[U|X] = E[U]$

- **Mean independence**. The **mean** of U is the same, regardless of the value of X :

These are **population moments**

- A **moment** is a specific attribute of a distribution
- The mean is the "first moment". Variance is the "second moment".

Economists spend a lot of time showing mean independence $E[U|X] = E[U]$.

Two quick reminders before we introduce the Ordinary Least Squares (OLS) estimator for β :

$$\text{Cov}(Y, X) = E[YX] - E[Y]E[X]$$

and

$$\text{If } E[U] = 0$$

then

$$\text{Cov}(U, X) = E[UX] - E[U]E[X] = E[UX] - 0$$

And note that the simple linear regression model $y = \beta_0 + \beta_1 x + u$ implies that:

$$u = y - \beta_0 - \beta_1 x$$

Since $u = y - \beta_0 - \beta_1 x$:

Let's write Assumption 1 and Assumption 2 using expectations of the **regression model from before**

- $E[U] = 0 \Rightarrow E[(y - \beta_0 - \beta_1 x)] = 0$
- $E[U|X] = 0 \Rightarrow E[x(y - \beta_0 - \beta_1 x)] = 0$
 - To see this, picture any expected value of x . Now, multiply it by 0.
- How many equations?
- How many unknowns?

Let's solve for β . To the board!

These are *moments*, and this way of deriving β is known as "method of moments".

What we just derived on the board depends on **population** moments: $Cov(Y, X)$ and $Var(Y, X)$.

But, just as before when we didn't know μ but we could calculate \bar{y} (and we even know something about the distribution of \bar{y})...

...we can calculate sample values for $Cov(y, x)$ and $Var(x)$

First, let's tackle the *estimate* of β_0 .

- We know, from the board, that $\beta_0 = E[y] - \beta_1 E[x]$
- We have a good, unbiased sample estimator for $E[y]$: \bar{y} .
- And we have a good, unbiased sample estimator for $E[X]$: \bar{x}
 - $\bar{y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{x}$

The hats stand for (sample) estimates! We don't observe β_0 , but we can estimate it. This is very common notation.

Of course, we still have to calculate $\hat{\beta}_1$.

We know how to calculate the sample covariance:

- $\widehat{Cov}(Y, X) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$


We know how to calculate the sample variance:

- $\widehat{Var}(X) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2$

$$\hat{\beta}_1 = \frac{\widehat{Cov}(Y, X)}{\widehat{Var}(X)} = \frac{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2}$$

What is important here is that **these are all observable in the data, and you know how to calculate them**. You know how to calculate \bar{x} and \bar{y} , you know how to sum things, and you know x_i and y_i in the data.

As long as your assumptions hold, you have an estimate of the PRF.



The red line is the *sample regression function*, or *SRF*.

Why is it the "sample" regression function?

A couple important terms:

- The **fitted value**, $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$
- The **residual**, $\hat{u}_i = y_i - \hat{y}_i = y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i$

And note that:

- $y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + \hat{u}_i$
 - The \hat{u}_i "trues up" the fitted value.

Note that the residual is not the same as the error term.

- The residual is an empirical estimate from the sample
- The error term, u_i , is different

What's inside the error term?


In u_i

- Omitted variables
 - There might be another covariate, x_2 , that is missing.
- Measurement error
 - That x might not be correctly measured.
- Non-linearities
 - Maybe there are some non-linear effects included in there.

These are all in u_i .

$$y_i = \beta_0 + \beta_1 x_1 + \underbrace{\beta_1(x_1^* - x_1) + \beta_2 x_{omitted} + f(nonlinears)}_{\text{other things, } u} + \tilde{u}_i$$


Our estimator, $\hat{\beta}$ assumes alllllll these things are 0 in expectation, no matter the value of x



Wooldridge Fig. 2.4

OLS in 1 variable

```
## `geom_smooth()` using formula 'y ~ x'
```



Regression line for wage2 data

It will always be the case that, for any estimates β from a sample:

- $\sum_{i=1}^N (\hat{u}_i) = 0$
- $\sum_{i=1}^N (x_i \hat{u}_i) = 0$
- The point (\bar{y}, \bar{x}) is always on the regression line

In a mathematical sense, we can always calculate a β such that $\bar{u} = 0$ for all values of x .

But what might throw us off is if there is something else unobserved, w , that is "in the error term".

- What is "in the error term?"
 - Everything in the world that isn't *educ*

For our example, let's think about $a = \textit{ability}$.

- Since *ability* is unobserved, it is "in the error term"
- *ability* is also pretty correlated with *educ* (high ability people go to college)
- In a way, we're attributing the causal effect of *ability* to *educ*
- So, $E[u | \textit{educ} = \textit{high}] > 0$
 - The unobserved error term, u , is higher due to the *ability* part

For now, we will work as if our assumption $E[U|X] = 0$ is true.

The "squares" part refers to the squaring of the error term.

The "least" part refers to a minimization of the (squared) error term.

Let's define the **sum of squared residuals** as:

$$SSR = \sum_{i=1}^N \hat{u}_i^2 = \sum_{i=1}^N (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2$$

And β is the "Least Squares" estimate if it minimizes SSR . How?

Take the derivative and set it equal to zero:

$$\frac{\partial SSR}{\partial \hat{\beta}_0} = 2 \sum (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) = 0$$

and

$$\frac{\partial SSR}{\partial \hat{\beta}_1} = 2 \sum (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) x_i = 0$$

SSR, SSE, and SST

We know that β_{OLS} minimizes the sum of squares. How do we measure how good of a fit we get?

Define two more in addition to SSR :

- Sum of Squares Total: $SST = \sum_{i=1}^N (y_i - \bar{y})^2$
 - SST is a total sum of squares (notice no hats).
- Sum of Squares Explained: $SSE = \sum_{i=1}^N (\hat{y}_i - \bar{y})^2$
 - SSE can be thought of as how much is *explained* by \hat{y}_i , ...relative to just *guessing the obvious*: \bar{y}

$$SST = SSR + SSE$$

The total variance is the sum of the variance of the residuals (what isn't explained by your model) and the *SSE* (the variance that is explained).

This is a *decomposition* of variance.

The R^2


R^2 ('r-square') is the comparison of SSE to SST . Since $SSE < SST$ always, and both are always positive, $0 < R^2 \leq 1$

$$R^2 = \frac{SSE}{SST} = 1 - \frac{SSR}{SST}$$

The R^2 is often interpreted as the "fraction of variance explained by the model"

- Your regression, the SRF, is a model
- The variance being explained is the variance in the outcome, y .

From earlier:



Wooldridge Fig. 2.4

```
# Ynum is the column name for the outcome variable  
# X is the column name for the independent variable and ex is the name of the  
summary(lm(Ynum ~ X, data=ex))
```

```
##  
## Call:  
## lm(formula = Ynum ~ X, data = ex)  
##  
## Residuals:  
##      1      2      3      4      5  
## -0.5046  0.5688 -0.8195  0.3760  0.3793  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)  -0.6530     0.6501  -1.004  0.38918  
## X              1.9591     0.2358   8.308  0.00365 **  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 0.7153 on 3 degrees of freedom  
## Multiple R-squared:  0.9583,    Adjusted R-squared:  0.9445  
## F-statistic: 69.02 on 1 and 3 DF,  p-value: 0.003654
```

Terminology

$$y = \beta_0 + \beta_1 x + u$$

y is called

- The dependent variable (DV)
- The "left hand side" (LHS)
- The outcome variable
- The response variable

x is called

- The independent variable
- The "right hand side" (RHS)
- The explanatory variable
- The control variable
- A covariate or a regressor

u is called

- The residual (when \hat{u})
- The error term (when u)

**On to
transformations
and functional
forms!**